

ReuseDistance

0.01

Generated by Doxygen 1.6.3

Tue Sep 25 11:06:42 2012

Contents

Chapter 1

Class Index

1.1 Class List

Here are the classes, structs, unions and interfaces with brief descriptions:

ReuseDistance	??
ReuseEntry	??
ReuseStats	??

Chapter 2

File Index

2.1 File List

Here is a list of all files with brief descriptions:

/home/michael/software/ReuseDistance/ReuseDistance.cpp	??
/home/michael/software/ReuseDistance/ReuseDistance.hpp	??

Chapter 3

Class Documentation

3.1 ReuseDistance Class Reference

```
#include <ReuseDistance.hpp>
```

Public Member Functions

- [ReuseDistance](#) ()
- [ReuseDistance](#) (uint64_t w)
- [ReuseDistance](#) ([ReuseDistance](#) &h)
- [~ReuseDistance](#) ()
- void [Print](#) ()
- void [Print](#) (std::ostream &f)
- void [Process](#) ([ReuseEntry](#) &addr)
- void [Process](#) ([ReuseEntry](#) *addrs, uint64_t count)
- void [Process](#) (std::vector< [ReuseEntry](#) > rs)
- void [Process](#) (std::vector< [ReuseEntry](#) * > addrs)
- uint64_t [GetDistance](#) ([ReuseEntry](#) &addr)
- [ReuseStats](#) * [GetStats](#) (uint64_t id)
- uint64_t [GetWindowSize](#) ()
- void [IncrementSequence](#) (uint64_t count)
- void [GetIndices](#) (std::vector< uint64_t > &ids)
- void [GetActiveAddresses](#) (std::vector< uint64_t > &addrs)
- uint64_t [GetSequenceValue](#) (uint64_t addr)
- uint64_t [GetCurrentSequence](#) ()
- void [Cleanup](#) ()
- void [SetCleanFrequency](#) (uint64_t c)
- uint64_t [GetCleanFrequency](#) ()

Static Public Attributes

- static const uint64_t [MinimumCleanFrequency](#) = 1000000

3.1.1 Detailed Description

Tracks reuse distances for a memory address stream. Keep track of the addresses within a specific window of history, whose size can be finite or infinite. See constructor documentation for more details.

Definition at line 185 of file ReuseDistance.hpp.

3.1.2 Constructor & Destructor Documentation

3.1.2.1 `ReuseDistance::ReuseDistance ()`

Constructs a [ReuseDistance](#) object. Default constructor. Uses an unlimited-size window.

Definition at line 25 of file ReuseDistance.cpp.

3.1.2.2 `ReuseDistance::ReuseDistance (uint64_t w)`

Constructs a [ReuseDistance](#) object.

Parameters

- w* The maximum size of the window of addresses that will be examined. Use 0 for no window, but be aware that this will use a potentially unlimited amount of memory that will be proportional to the number of unique addresses processed by this object.

Definition at line 33 of file ReuseDistance.cpp.

3.1.2.3 `ReuseDistance::ReuseDistance (ReuseDistance & h)`

Constructs a [ReuseDistance](#) object. Copy constructor. Performs a deep copy.

Parameters

- h* A reference to another [ReuseDistance](#) object. All state from this parameter is copied to the new [ReuseDistance](#) object, including window size, current addresses in that window and all tracked statistics.

Definition at line 41 of file ReuseDistance.cpp.

3.1.2.4 `ReuseDistance::~~ReuseDistance ()`

Destroys a [ReuseDistance](#) object.

Definition at line 72 of file ReuseDistance.cpp.

3.1.3 Member Function Documentation

3.1.3.1 `void ReuseDistance::Cleanup ()`

Clean up the address window. That is, reclaim memory for all addresses strictly outside the maximum window size. This is done for you periodically so you don't ever really need to call this yourself.

Returns

none

Definition at line 121 of file ReuseDistance.cpp.

3.1.3.2 void ReuseDistance::GetActiveAddresses (std::vector< uint64_t > & *addrs*)

Get a std::vector containing all of the addresses currently in this [ReuseDistance](#) object's active window.

Parameters

addrs A std::vector which will contain the addresses. It is an error to pass this vector non-empty (that is `addrs.size() == 0` is enforced).

Returns

none

Definition at line 87 of file ReuseDistance.cpp.

3.1.3.3 uint64_t ReuseDistance::GetCleanFrequency () [inline]

Get the frequency with which Clean is called. Has no meaning when window size is unlimited.

Returns

The frequency with which Cleanup is called.

Definition at line 408 of file ReuseDistance.hpp.

3.1.3.4 uint64_t ReuseDistance::GetCurrentSequence () [inline]

Get this [ReuseDistance](#) object's current sequence.

Returns

This [ReuseDistance](#) object's current sequence.

Definition at line 380 of file ReuseDistance.hpp.

3.1.3.5 uint64_t ReuseDistance::GetDistance (ReuseEntry & *addr*)

Get a reuse distance for a memory address without tracking statistics for it.

Parameters

addr The memory address to analyze.

Returns

The reuse distance for the memory address given by *addr*.

3.1.3.6 void ReuseDistance::GetIndices (std::vector< uint64_t > & *ids*)

Get a std::vector containing all of the unique indices processed by this [ReuseDistance](#) object.

Parameters

ids A std::vector which will contain the ids. It is an error to pass this vector non-empty (that is `ids.size() == 0` is enforced).

Returns

none

Definition at line 79 of file ReuseDistance.cpp.

3.1.3.7 uint64_t ReuseDistance::GetSequenceValue (uint64_t *addr*)

Get the sequence value for an address currently in this [ReuseDistance](#) object's active window.

Parameters

addr An address. Addresses not in this object's active window will generate a return value of 0.

Returns

The sequence value for *addr*, or 0 if *addr* is not in this object's active window.

Definition at line 95 of file ReuseDistance.cpp.

3.1.3.8 ReuseStats * ReuseDistance::GetStats (uint64_t *id*)

Get the [ReuseStats](#) object associated with some unique id.

Parameters

id The unique id.

Returns

The [ReuseStats](#) object associated with parameter *id*.

Definition at line 195 of file ReuseDistance.cpp.

3.1.3.9 uint64_t ReuseDistance::GetWindowSize () [inline]

Get the size of the window for this [ReuseDistance](#) object.

Returns

The size of the window for this [ReuseDistance](#) object.

Definition at line 328 of file ReuseDistance.hpp.

3.1.3.10 void ReuseDistance::IncrementSequence (uint64_t *count*) [inline]

Increment the internal sequence count for this [ReuseDistance](#) object. This has the effect of fast forwarding in the memory address stream. Possibly useful if you are using sampling on your memory address stream.

Parameters

count The amount of the increment.

Returns

none

Definition at line 339 of file ReuseDistance.hpp.

3.1.3.11 void ReuseDistance::Print (std::ostream & *f*)

Print statistics for this [ReuseDistance](#) to an output stream. The first line of the output is five tokens which are [1] the string literal REUSESTATS, [2] the unique id, [3] the window size (0 == unlimited) [4] the total number of accesses for that unique id and [5] the number of accesses from that id which were not found within the active address window either because they were evicted or because of cold misses. Each additional line of output contains two tokens, which give [1] a reuse distance and [2] the number of times that reuse distance was observed.

Parameters

f The output stream to print results to.

Returns

none

3.1.3.12 void ReuseDistance::Print ()

Print statistics for this [ReuseDistance](#) to std::cout. See the other version of [ReuseDistance::Print](#) for information about output format.

Returns

none

Definition at line 103 of file ReuseDistance.cpp.

3.1.3.13 void ReuseDistance::Process (std::vector< ReuseEntry * > *addrs*)

Process multiple memory addresses.

Parameters

addrs A std::vector of memory addresses to process.

Returns

none

3.1.3.14 void ReuseDistance::Process (std::vector< ReuseEntry > rs)

Process multiple memory addresses.

Parameters

addr A std::vector of memory addresses to process.

Returns

none

3.1.3.15 void ReuseDistance::Process (ReuseEntry * addr, uint64_t count)

Process multiple memory addresses.

Parameters

addr An array of structures describing memory addresses to process.

count The number of elements in addr.

Returns

none

Definition at line 148 of file ReuseDistance.cpp.

3.1.3.16 void ReuseDistance::Process (ReuseEntry & addr)

Process a single memory address.

Parameters

addr The structure describing the memory address to process.

Returns

none

Definition at line 168 of file ReuseDistance.cpp.

3.1.3.17 void ReuseDistance::SetCleanFrequency (uint64_t c)

Set the frequency with which Clean is called. By default this is defined using the minimum of [ReuseDistance::MinimumCleanFrequency](#) and the window size. Has no meaning when window size is unlimited.

Parameters

c The frequency with which to call Cleanup.

Returns

none

Definition at line 67 of file ReuseDistance.cpp.

3.1.4 Member Data Documentation

3.1.4.1 `const uint64_t ReuseDistance::MinimumCleanFrequency = 1000000` `[static]`

Minimum value for the cleanup frequency. The cleanup frequency is set by the constructor to the maximum of this value and the window size.

Definition at line 212 of file ReuseDistance.hpp.

The documentation for this class was generated from the following files:

- [/home/michaell/software/ReuseDistance/ReuseDistance.hpp](#)
- [/home/michaell/software/ReuseDistance/ReuseDistance.cpp](#)

3.2 ReuseEntry Struct Reference

```
#include <ReuseDistance.hpp>
```

Public Attributes

- uint64_t [id](#)
- uint64_t [address](#)

3.2.1 Detailed Description

[ReuseEntry](#) is used to pass memory addresses into a [ReuseDistance](#).

id The unique id of the entity which generated the memory address. Statistics are tracked separately for each unique id. **address** A memory address.

Definition at line 62 of file [ReuseDistance.hpp](#).

3.2.2 Member Data Documentation

3.2.2.1 uint64_t ReuseEntry::address

Definition at line 64 of file [ReuseDistance.hpp](#).

3.2.2.2 uint64_t ReuseEntry::id

Definition at line 63 of file [ReuseDistance.hpp](#).

The documentation for this struct was generated from the following file:

- [/home/michaell/software/ReuseDistance/ReuseDistance.hpp](#)

3.3 ReuseStats Class Reference

```
#include <ReuseDistance.hpp>
```

Public Member Functions

- [ReuseStats](#) ()
- [ReuseStats](#) ([ReuseStats](#) &r)
- [~ReuseStats](#) ()
- void [Update](#) (uint64_t dist)
- void [Miss](#) ()
- uint64_t [GetMissCount](#) ()
- void [Print](#) (std::ostream &f, uint64_t uniqueid)
- void [GetSortedDistances](#) (std::vector< uint64_t > &dists)
- uint64_t [GetMaximumDistance](#) ()
- uint64_t [CountDistance](#) (uint64_t dist)
- uint64_t [CountDistance](#) (uint64_t low, uint64_t high)
- uint64_t [GetAccessCount](#) ()

3.3.1 Detailed Description

[ReuseStats](#) holds a count of the number of times each reuse distance is observed.

Definition at line 72 of file [ReuseDistance.hpp](#).

3.3.2 Constructor & Destructor Documentation

3.3.2.1 [ReuseStats::ReuseStats](#) () `[inline]`

Constructs a [ReuseStats](#) object. Default constructor.

Definition at line 83 of file [ReuseDistance.hpp](#).

3.3.2.2 [ReuseStats::ReuseStats](#) ([ReuseStats](#) & r)

Constructs a [ReuseStats](#) object. Copy constructor. Performs a deep copy.

Parameters

r A [ReuseStats](#) object from which to copy all state.

Definition at line 205 of file [ReuseDistance.cpp](#).

3.3.2.3 [ReuseStats::~~ReuseStats](#) () `[inline]`

Destroys a [ReuseStats](#) object.

Definition at line 95 of file [ReuseDistance.hpp](#).

3.3.3 Member Function Documentation

3.3.3.1 `uint64_t ReuseStats::CountDistance (uint64_t low, uint64_t high)`

Count the number of times any distance within some range [*low*, *high*) has been observed.

Parameters

low The lower bound (inclusive) of the distance range to count.

high The upper bound (exclusive) of the distance range to count.

Returns

The number of times any distance within the range [*low*, *high*) has been observed.

Definition at line 250 of file ReuseDistance.cpp.

3.3.3.2 `uint64_t ReuseStats::CountDistance (uint64_t dist)`

Count the number of times some distance has been observed.

Parameters

dist The distance to count.

Returns

The number of times *d* has been observed.

Definition at line 243 of file ReuseDistance.cpp.

3.3.3.3 `uint64_t ReuseStats::GetAccessCount ()`

Count the total number of distances observed.

Returns

The total number of distances observed.

Definition at line 218 of file ReuseDistance.cpp.

3.3.3.4 `uint64_t ReuseStats::GetMaximumDistance ()`

Get the maximum distance observed.

Returns

The maximum distance observed.

Definition at line 222 of file ReuseDistance.cpp.

3.3.3.5 uint64_t ReuseStats::GetMissCount () [inline]

Get the number of misses. This is equal to the number of times [Miss\(\)](#) is called plus the number of times [Update\(0\)](#) is called. These two calls are functionally equivalent but [Miss\(\)](#) is faster.

Definition at line 119 of file [ReuseDistance.hpp](#).

3.3.3.6 void ReuseStats::GetSortedDistances (std::vector< uint64_t > & dists)

Get a `std::vector` containing the distances observed, sorted in ascending order.

Parameters

dists The vector which will hold the sorted distance values. It is an error for *dists* to be passed in non-empty (that is, `dists.size() == 0` is enforced).

Returns

none

3.3.3.7 void ReuseStats::Miss ()

Increment the number of misses. That is, addresses which were not found inside the active address window. This is equivalent [Update\(0\)](#), but is faster.

Returns

none

Definition at line 238 of file [ReuseDistance.cpp](#).

3.3.3.8 void ReuseStats::Print (std::ostream &f, uint64_t uniqueid)

Print a summary of the current reuse distances and counts.

Parameters

f The stream to receive the output.

uniqueid An identifier for this [ReuseStats](#) object.

scale A vector holding the boundaries of bins used to aggregate the reuse distances.

Returns

none

3.3.3.9 void ReuseStats::Update (uint64_t dist)

Increment the counter for some distance.

Parameters

dist A reuse distance observed in the memory address stream.

Returns

none

Definition at line 233 of file ReuseDistance.cpp.

The documentation for this class was generated from the following files:

- </home/michaell/software/ReuseDistance/ReuseDistance.hpp>
- </home/michaell/software/ReuseDistance/ReuseDistance.cpp>

Chapter 4

File Documentation

4.1 /home/michaell/software/ReuseDistance/ReuseDistance.cpp File Reference

```
#include <ReuseDistance.hpp>
```

4.2 /home/michaell/software/ReuseDistance/ReuseDistance.hpp File Reference

```
#include <assert.h>
#include <stdlib.h>
#include <algorithm>
#include <iostream>
#include <ostream>
#include <set>
#include <vector>
#include <map>
```

Classes

- struct [ReuseEntry](#)
- class [ReuseStats](#)
- class [ReuseDistance](#)

Defines

- #define [reuse_map_type](#) std::map
- #define [TAB](#) "\\t"
- #define [ENDL](#) "\\n"

4.2.1 Detailed Description

Author

Michael Laurenzano <michaell@sdsc.edu>

Version

0.01

4.2.2 LICENSE

This file is part of the [ReuseDistance](#) tool.

Copyright (c) 2012, University of California Regents All rights reserved.

This program is free software: you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, either version 3 of the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details.

You should have received a copy of the GNU General Public License along with this program. If not, see <<http://www.gnu.org/licenses/>>.

4.2.3 DESCRIPTION

The ReuseDistanceHandler class allows for calculation and statistic tracking for finding memory reuse distances given a stream of memory addresses and ids.

Definition in file [ReuseDistance.hpp](#).

4.2.4 Define Documentation

4.2.4.1 `#define ENDL "\n"`

Definition at line 51 of file ReuseDistance.hpp.

4.2.4.2 `#define reuse_map_type std::map`

Definition at line 47 of file ReuseDistance.hpp.

4.2.4.3 `#define TAB "\t"`

Definition at line 50 of file ReuseDistance.hpp.